

```
!pip install torch torchvision matplotlib seaborn scikit-learn
```

```
Requirement already satisfied: torch in
/usr/local/lib/python3.11/dist-packages (2.5.1+cu121)
Requirement already satisfied: torchvision in
/usr/local/lib/python3.11/dist-packages (0.20.1+cu121)
Requirement already satisfied: matplotlib in
/usr/local/lib/python3.11/dist-packages (3.10.0)
Requirement already satisfied: seaborn in
/usr/local/lib/python3.11/dist-packages (0.13.2)
Requirement already satisfied: scikit-learn in
/usr/local/lib/python3.11/dist-packages (1.6.1)
Requirement already satisfied: filelock in
/usr/local/lib/python3.11/dist-packages (from torch) (3.17.0)
Requirement already satisfied: typing-extensions>=4.8.0 in
/usr/local/lib/python3.11/dist-packages (from torch) (4.12.2)
Requirement already satisfied: networkx in
/usr/local/lib/python3.11/dist-packages (from torch) (3.4.2)
Requirement already satisfied: jinja2 in
/usr/local/lib/python3.11/dist-packages (from torch) (3.1.5)
Requirement already satisfied: fsspec in
/usr/local/lib/python3.11/dist-packages (from torch) (2024.10.0)
Requirement already satisfied: nvidia-cuda-nvrtc-cu12==12.1.105 in
/usr/local/lib/python3.11/dist-packages (from torch) (12.1.105)
Requirement already satisfied: nvidia-cuda-runtime-cu12==12.1.105
in /usr/local/lib/python3.11/dist-packages (from torch) (12.1.105)
Requirement already satisfied: nvidia-cuda-cupti-cu12==12.1.105 in
/usr/local/lib/python3.11/dist-packages (from torch) (12.1.105)
Requirement already satisfied: nvidia-cudnn-cu12==9.1.0.70 in
/usr/local/lib/python3.11/dist-packages (from torch) (9.1.0.70)
Requirement already satisfied: nvidia-cublas-cu12==12.1.3.1 in
/usr/local/lib/python3.11/dist-packages (from torch) (12.1.3.1)
Requirement already satisfied: nvidia-cufft-cu12==11.0.2.54 in
/usr/local/lib/python3.11/dist-packages (from torch) (11.0.2.54)
Requirement already satisfied: nvidia-curand-cu12==10.3.2.106 in
/usr/local/lib/python3.11/dist-packages (from torch) (10.3.2.106)
Requirement already satisfied: nvidia-cusolver-cu12==11.4.5.107 in
/usr/local/lib/python3.11/dist-packages (from torch) (11.4.5.107)
Requirement already satisfied: nvidia-cuspars-cu12==12.1.0.106 in
/usr/local/lib/python3.11/dist-packages (from torch) (12.1.0.106)
Requirement already satisfied: nvidia-nccl-cu12==2.21.5 in
/usr/local/lib/python3.11/dist-packages (from torch) (2.21.5)
Requirement already satisfied: nvidia-nvtx-cu12==12.1.105 in
/usr/local/lib/python3.11/dist-packages (from torch) (12.1.105)
Requirement already satisfied: triton==3.1.0 in
/usr/local/lib/python3.11/dist-packages (from torch) (3.1.0)
Requirement already satisfied: sympy==1.13.1 in
/usr/local/lib/python3.11/dist-packages (from torch) (1.13.1)
Requirement already satisfied: nvidia-nvjitlink-cu12 in
/usr/local/lib/python3.11/dist-packages (from nvidia-cusolver-
```

```
cul2==11.4.5.107->torch) (12.6.85)
Requirement already satisfied: mpmath<1.4,>=1.1.0 in
/usr/local/lib/python3.11/dist-packages (from sympy==1.13.1->torch)
(1.3.0)
Requirement already satisfied: numpy in
/usr/local/lib/python3.11/dist-packages (from torchvision) (1.26.4)
Requirement already satisfied: pillow!=8.3.*,>=5.3.0 in
/usr/local/lib/python3.11/dist-packages (from torchvision) (11.1.0)
Requirement already satisfied: contourpy>=1.0.1 in
/usr/local/lib/python3.11/dist-packages (from matplotlib) (1.3.1)
Requirement already satisfied: cycler>=0.10 in
/usr/local/lib/python3.11/dist-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in
/usr/local/lib/python3.11/dist-packages (from matplotlib) (4.55.4)
Requirement already satisfied: kiwisolver>=1.3.1 in
/usr/local/lib/python3.11/dist-packages (from matplotlib) (1.4.8)
Requirement already satisfied: packaging>=20.0 in
/usr/local/lib/python3.11/dist-packages (from matplotlib) (24.2)
Requirement already satisfied: pyparsing>=2.3.1 in
/usr/local/lib/python3.11/dist-packages (from matplotlib) (3.2.1)
Requirement already satisfied: python-dateutil>=2.7 in
/usr/local/lib/python3.11/dist-packages (from matplotlib) (2.8.2)
Requirement already satisfied: pandas>=1.2 in
/usr/local/lib/python3.11/dist-packages (from seaborn) (2.2.2)
Requirement already satisfied: scipy>=1.6.0 in
/usr/local/lib/python3.11/dist-packages (from scikit-learn) (1.13.1)
Requirement already satisfied: joblib>=1.2.0 in
/usr/local/lib/python3.11/dist-packages (from scikit-learn) (1.4.2)
Requirement already satisfied: threadpoolctl>=3.1.0 in
/usr/local/lib/python3.11/dist-packages (from scikit-learn) (3.5.0)
Requirement already satisfied: pytz>=2020.1 in
/usr/local/lib/python3.11/dist-packages (from pandas>=1.2->seaborn)
(2024.2)
Requirement already satisfied: tzdata>=2022.7 in
/usr/local/lib/python3.11/dist-packages (from pandas>=1.2->seaborn)
(2025.1)
Requirement already satisfied: six>=1.5 in
/usr/local/lib/python3.11/dist-packages (from python-dateutil>=2.7-
>matplotlib) (1.17.0)
Requirement already satisfied: MarkupSafe>=2.0 in
/usr/local/lib/python3.11/dist-packages (from jinja2->torch) (3.0.2)
```

```
import torch
import torch.nn as nn
import torch.optim as optim
from torchvision import datasets, transforms, models
import os
from sklearn.metrics import confusion_matrix, classification_report
import numpy as np
import matplotlib.pyplot as plt
```

```

from PIL import Image
import seaborn as sns
from torch.optim import lr_scheduler

# Define data transformations for data augmentation and normalization
data_transforms = {
    'train': transforms.Compose([
        transforms.RandomResizedCrop(224),
        transforms.RandomHorizontalFlip(),
        transforms.ColorJitter(brightness=0.2, contrast=0.2,
saturation=0.2, hue=0.2),
        transforms.RandomRotation(30),
        transforms.ToTensor(),
        transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224,
0.225])
    ]),
    'test': transforms.Compose([ # Changed to test dataset instead of
val
        transforms.Resize(256),
        transforms.CenterCrop(224),
        transforms.ToTensor(),
        transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224,
0.225])
    ]),
}

# Define the data directory
data_dir = '/content/drive/MyDrive/Images'

# Create data loaders
image_datasets = {x: datasets.ImageFolder(os.path.join(data_dir, x),
data_transforms[x]) for x in ['train', 'test']}
dataloaders = {x: torch.utils.data.DataLoader(image_datasets[x],
batch_size=16, shuffle=True, num_workers=4) for x in ['train',
'test']}
dataset_sizes = {x: len(image_datasets[x]) for x in ['train', 'test']}
print(dataset_sizes)

class_names = image_datasets['train'].classes

{'train': 3902, 'test': 977}

/usr/local/lib/python3.11/dist-packages/torch/utils/data/
dataloader.py:617: UserWarning: This DataLoader will create 4 worker
processes in total. Our suggested max number of worker in current
system is 2, which is smaller than what this DataLoader is going to
create. Please be aware that excessive worker creation might get
DataLoader running slow or even freeze, lower the worker number to
avoid potential slowness/freeze if necessary.
  warnings.warn(

```

```

# Load the pre-trained ResNet-50 model
model = models.resnet50(pretrained=True)

# Freeze all layers except the final classification layer
for name, param in model.named_parameters():
    if "fc" in name: # Unfreeze the final classification layer
        param.requires_grad = True
    else:
        param.requires_grad = False

# Modify the final fully connected layer for 6 classes
model.fc = nn.Linear(model.fc.in_features, 6)

/usr/local/lib/python3.11/dist-packages/torchvision/models/_
_utils.py:208: UserWarning: The parameter 'pretrained' is deprecated
since 0.13 and may be removed in the future, please use 'weights'
instead.
  warnings.warn(
/usr/local/lib/python3.11/dist-packages/torchvision/models/_utils.py:23: UserWarning: Arguments other than a weight enum or `None` for
'weights' are deprecated since 0.13 and may be removed in the future.
The current behavior is equivalent to passing
`weights=ResNet50_Weights.IMAGENET1K_V1`. You can also use
`weights=ResNet50_Weights.DEFAULT` to get the most up-to-date weights.
  warnings.warn(msg)
Downloading: "https://download.pytorch.org/models/resnet50-
0676ba61.pth" to /root/.cache/torch/hub/checkpoints/resnet50-
0676ba61.pth
100%|██████████| 97.8M/97.8M [00:00<00:00, 140MB/s]

# Define the loss function and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(model.parameters(), lr=0.001, momentum=0.9,
weight_decay=0.0001) # Weight decay added

# Learning Rate Scheduler
scheduler = lr_scheduler.StepLR(optimizer, step_size=7, gamma=0.1)

# Move the model to the GPU if available
device = torch.device("cuda:0" if torch.cuda.is_available() else
"cpu")
model = model.to(device)

# Training loop with early stopping
num_epochs = 15
best_model_wts = model.state_dict()
best_acc = 0.0
train_losses, test_losses = [], []
train_accs, test_accs = [], []

```

```

for epoch in range(num_epochs):
    print(f'Epoch {epoch+1}/{num_epochs}')
    print('-' * 10)

    for phase in ['train', 'test']:
        if phase == 'train':
            model.train()
        else:
            model.eval()

        running_loss = 0.0
        running_corrects = 0

        for inputs, labels in dataloaders[phase]:
            inputs = inputs.to(device)
            labels = labels.to(device)

            optimizer.zero_grad()

            with torch.set_grad_enabled(phase == 'train'):
                outputs = model(inputs)
                _, preds = torch.max(outputs, 1)
                loss = criterion(outputs, labels)

                if phase == 'train':
                    loss.backward()
                    optimizer.step()

            running_loss += loss.item() * inputs.size(0)
            running_corrects += torch.sum(preds == labels.data)

        epoch_loss = running_loss / dataset_sizes[phase]
        epoch_acc = running_corrects.double() / dataset_sizes[phase]
        if phase == 'train':
            train_losses.append(epoch_loss)
            train_accs.append(epoch_acc)
        else:
            test_losses.append(epoch_loss)
            test_accs.append(epoch_acc)

    print(f'{phase} Loss: {epoch_loss:.4f} Acc: {epoch_acc:.4f}')

    # Deep copy the model
    if phase == 'test' and epoch_acc > best_acc:
        best_acc = epoch_acc
        best_model_wts = model.state_dict()

    scheduler.step() # Update the learning rate

```

```
# Load best model weights
model.load_state_dict(best_model_wts)
```

```
Epoch 1/15
```

```
-----
train Loss: 1.1177 Acc: 0.6184
test Loss: 0.5729 Acc: 0.8362
```

```
Epoch 2/15
```

```
-----
train Loss: 0.7195 Acc: 0.7629
test Loss: 0.4564 Acc: 0.8547
```

```
Epoch 3/15
```

```
-----
train Loss: 0.6553 Acc: 0.7750
test Loss: 0.3772 Acc: 0.8772
```

```
Epoch 4/15
```

```
-----
train Loss: 0.6080 Acc: 0.7916
test Loss: 0.3373 Acc: 0.8895
```

```
Epoch 5/15
```

```
-----
train Loss: 0.5934 Acc: 0.7932
test Loss: 0.2980 Acc: 0.9079
```

```
Epoch 6/15
```

```
-----
train Loss: 0.5817 Acc: 0.7996
test Loss: 0.3024 Acc: 0.8915
```

```
Epoch 7/15
```

```
-----
train Loss: 0.5763 Acc: 0.8063
test Loss: 0.3243 Acc: 0.8802
```

```
Epoch 8/15
```

```
-----
train Loss: 0.5085 Acc: 0.8234
test Loss: 0.2847 Acc: 0.9079
```

```
Epoch 9/15
```

```
-----
train Loss: 0.5200 Acc: 0.8255
test Loss: 0.2873 Acc: 0.9089
```

```
Epoch 10/15
```

```
-----
train Loss: 0.5060 Acc: 0.8270
test Loss: 0.2977 Acc: 0.8946
```

```
Epoch 11/15
```

```
-----
train Loss: 0.5219 Acc: 0.8232
test Loss: 0.2800 Acc: 0.9120
```

```
Epoch 12/15
```

```
-----
train Loss: 0.5295 Acc: 0.8209
```

```
test Loss: 0.2717 Acc: 0.9161
Epoch 13/15
```

```
-----
```

```
train Loss: 0.5183 Acc: 0.8203
test Loss: 0.2897 Acc: 0.9017
Epoch 14/15
```

```
-----
```

```
train Loss: 0.5069 Acc: 0.8247
test Loss: 0.2770 Acc: 0.9079
Epoch 15/15
```

```
-----
```

```
train Loss: 0.5027 Acc: 0.8314
test Loss: 0.3123 Acc: 0.8895
```

```
<All keys matched successfully>
```

```
# Save the trained model
```

```
torch.save(model.state_dict(),
'best_classification_model_resnet50.pth')
```

```
# Evaluate model performance
```

```
def evaluate_model(model, dataloader):
```

```
    model.eval()
```

```
    all_preds = []
```

```
    all_labels = []
```

```
    with torch.no_grad():
```

```
        for inputs, labels in dataloaders['test']: # Using test set
for evaluation
```

```
            inputs = inputs.to(device)
```

```
            labels = labels.to(device)
```

```
            outputs = model(inputs)
```

```
            _, preds = torch.max(outputs, 1)
```

```
            all_preds.extend(preds.cpu().numpy())
```

```
            all_labels.extend(labels.cpu().numpy())
```

```
    cm = confusion_matrix(all_labels, all_preds)
```

```
    cr = classification_report(all_labels, all_preds,
target_names=class_names)
```

```
    return cm, cr
```

```
# Get confusion matrix and classification report
```

```
cm, cr = evaluate_model(model, dataloaders['test'])
```


```
print("Classification Report:\n", cr)
```

```
Classification Report:
```

	precision	recall	f1-score	support
Aphids	0.81	0.95	0.87	164
Army worm	0.90	0.93	0.92	162
Bacterial blight	0.95	0.81	0.88	160

Healthy	0.85	0.99	0.91	165
Powdery mildew	0.98	0.79	0.87	166
Target spot	0.90	0.87	0.89	160
accuracy			0.89	977
macro avg	0.90	0.89	0.89	977
weighted avg	0.90	0.89	0.89	977

```
# Visualize Confusion Matrix
plt.figure(figsize=(10, 7))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=class_names, yticklabels=class_names)
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()
```



```
# Example image prediction with probabilities
image_path = '/content/powdery-mildew-winter-squash.jpg' # Replace
with your test image path
image = Image.open(image_path)
preprocess = transforms.Compose([
```



```

        transforms.Resize(256),
        transforms.CenterCrop(224),
        transforms.ToTensor(),
        transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225]))
    ])
    input_tensor = preprocess(image)
    input_batch = input_tensor.unsqueeze(0).to(device)

    # Perform inference
    model.eval()
    with torch.no_grad():
        output = model(input_batch)
        probabilities = torch.nn.functional.softmax(output, dim=1)
        confidence, predicted_class = torch.max(probabilities, 1)

    # Get predicted class and probability
    predicted_class_name = class_names[predicted_class.item()]
    predicted_prob = probabilities[0][predicted_class].item()


    print(f'The predicted class is: {predicted_class_name} with a
    confidence of {predicted_prob:.4f}')

    # Show the image with prediction
    plt.imshow(np.array(image))
    plt.axis('off')
    plt.text(10, 10, f'Predicted: {predicted_class_name}
    ({predicted_prob*100:.2f}%)', fontsize=12, color='white',
    backgroundColor='red')
    plt.show()

```

The predicted class is: Powdery mildew with a confidence of 0.9562

Predicted: Powdery mildew (95.62%)



```
# Ensure all tensors in train_accs and test_accs are converted to
Python floats
train_accs = [acc.cpu().item() if isinstance(acc, torch.Tensor) else
acc for acc in train_accs]
test_accs = [acc.cpu().item() if isinstance(acc, torch.Tensor) else
acc for acc in test_accs]
```

```
# Plotting
```

```
epochs = range(1, num_epochs + 1)
plt.figure(figsize=(12, 5))
```

```
# Loss plot
```

```
plt.subplot(1, 2, 1)
plt.plot(epochs, train_losses, label='Train Loss')
plt.plot(epochs, test_losses, label='Test Loss')
plt.title('Loss Curve')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
```

```
# Accuracy plot
```

```
plt.subplot(1, 2, 2)
plt.plot(epochs, train_accs, label='Train Accuracy')
plt.plot(epochs, test_accs, label='Test Accuracy')
plt.title('Accuracy Curve')
plt.xlabel('Epochs')
```

```
plt.ylabel('Accuracy')  
plt.legend()
```

```
plt.tight_layout()  
plt.show()
```

